

Streaming Data Engineering

A Tale of Four Streams

Emanuele Della Valle

prof @ Politecnico di Milano

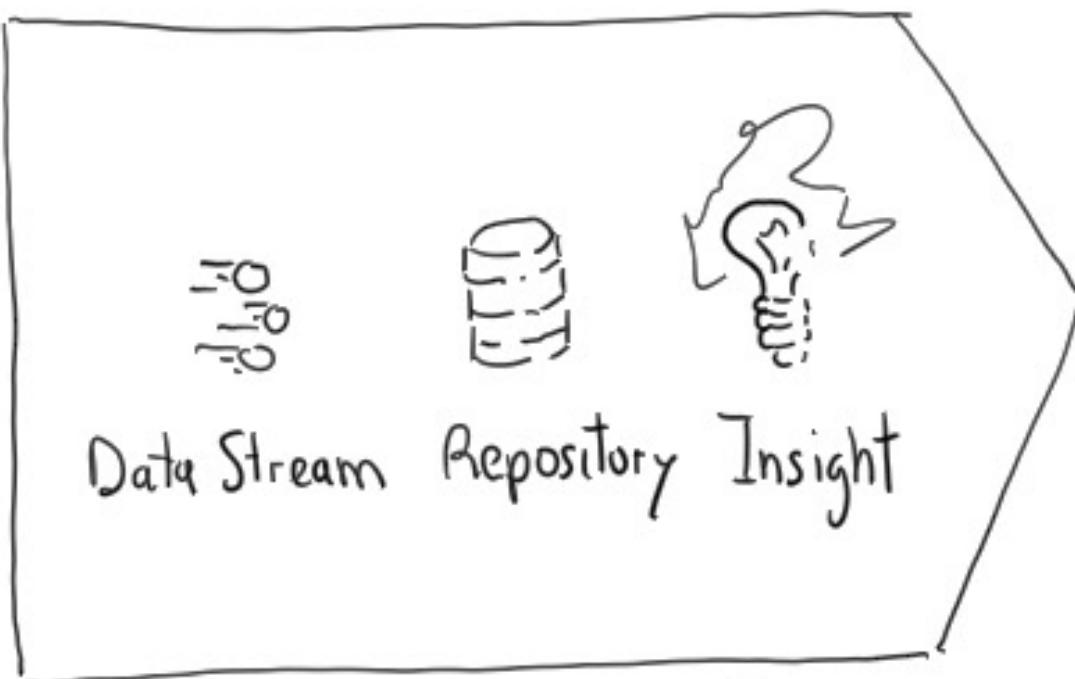
founder @ Quantia Consulting

1st premises: continuity matters



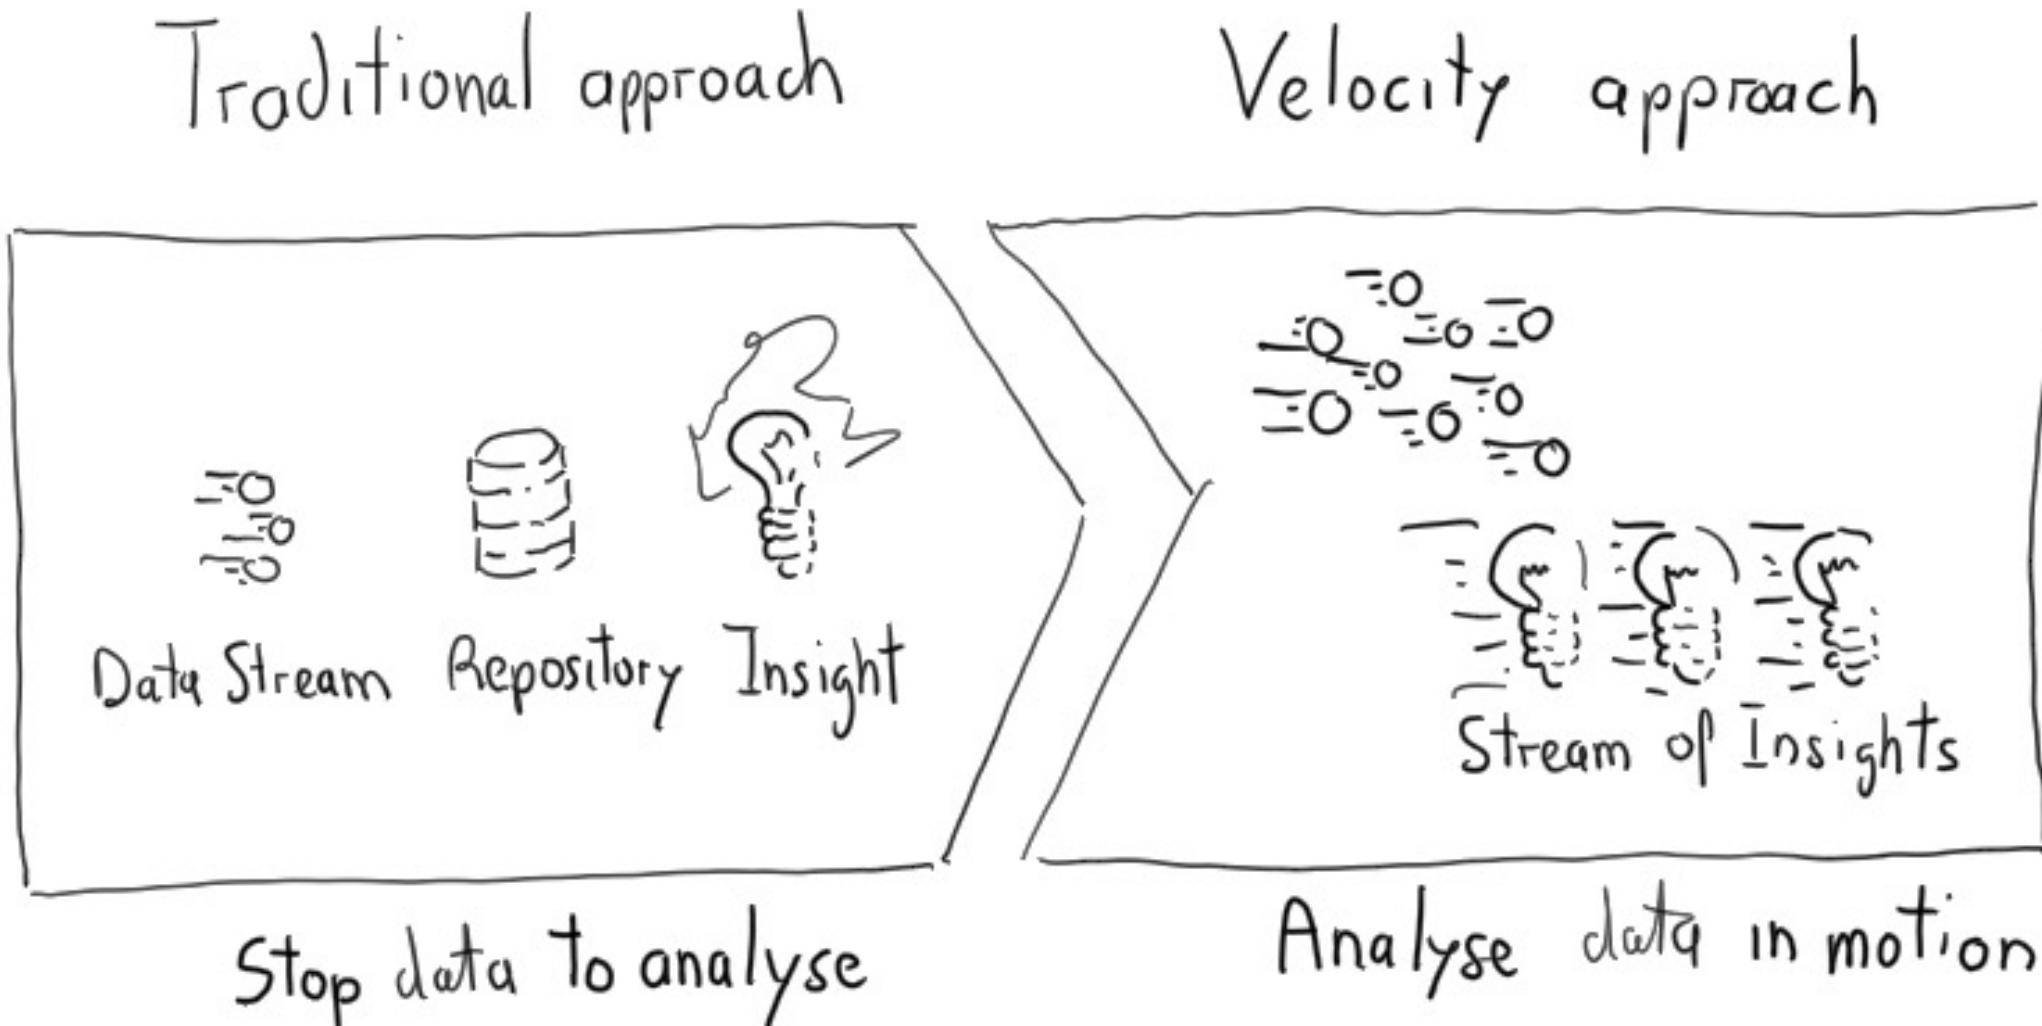
1st premises: continuity matters

Traditional approach



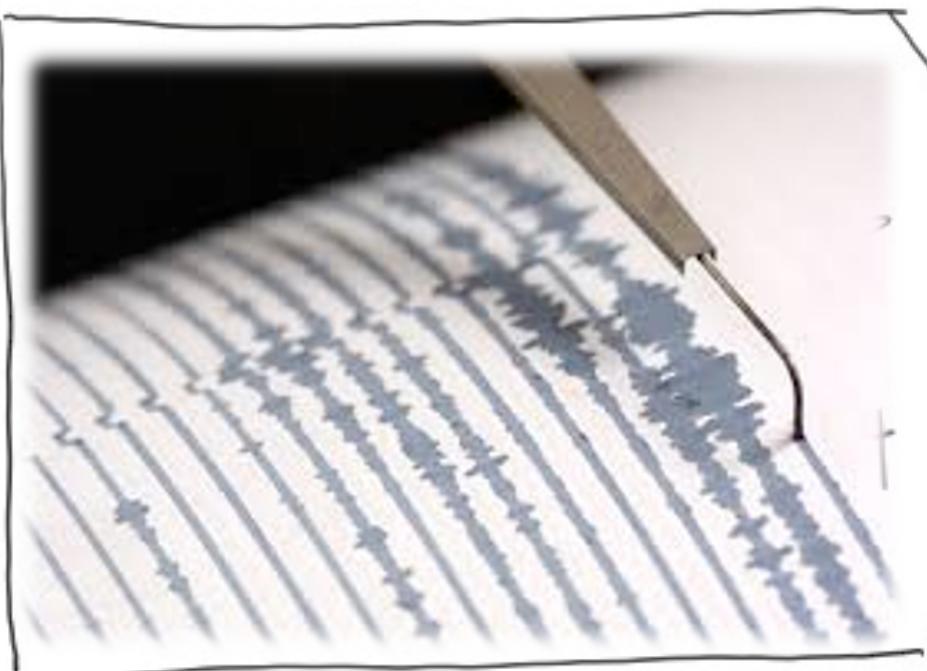
Stop data to analyse

1st premises: continuity matters



1st premises: continuity matters

Traditional approach



Stop data to analyse

Velocity approach

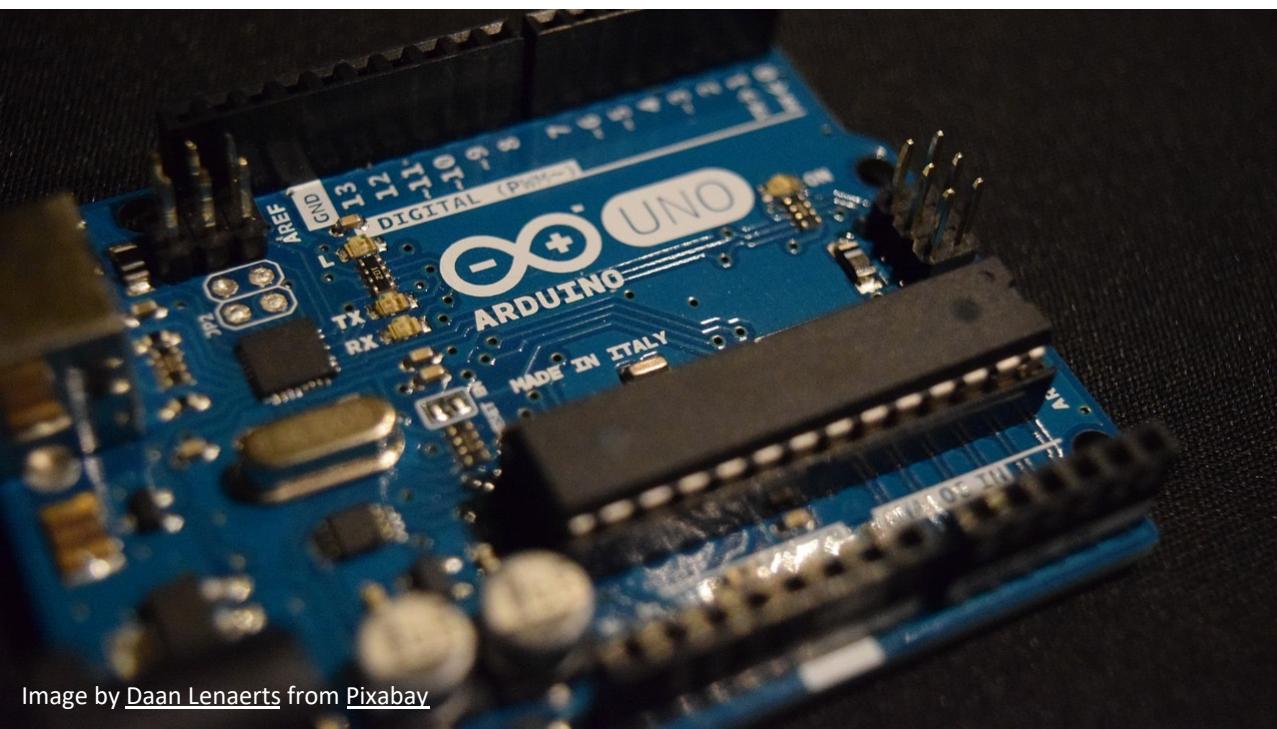


Analyse data in motion

**Stream: a useful metaphor to
explore the solution space**

Streams take many forms ...

**Myriads of tiny flows
that you can collect**



**e.g., from physical or
software sensors**



e.g., social media,
telcos' and utilities'
monitoring

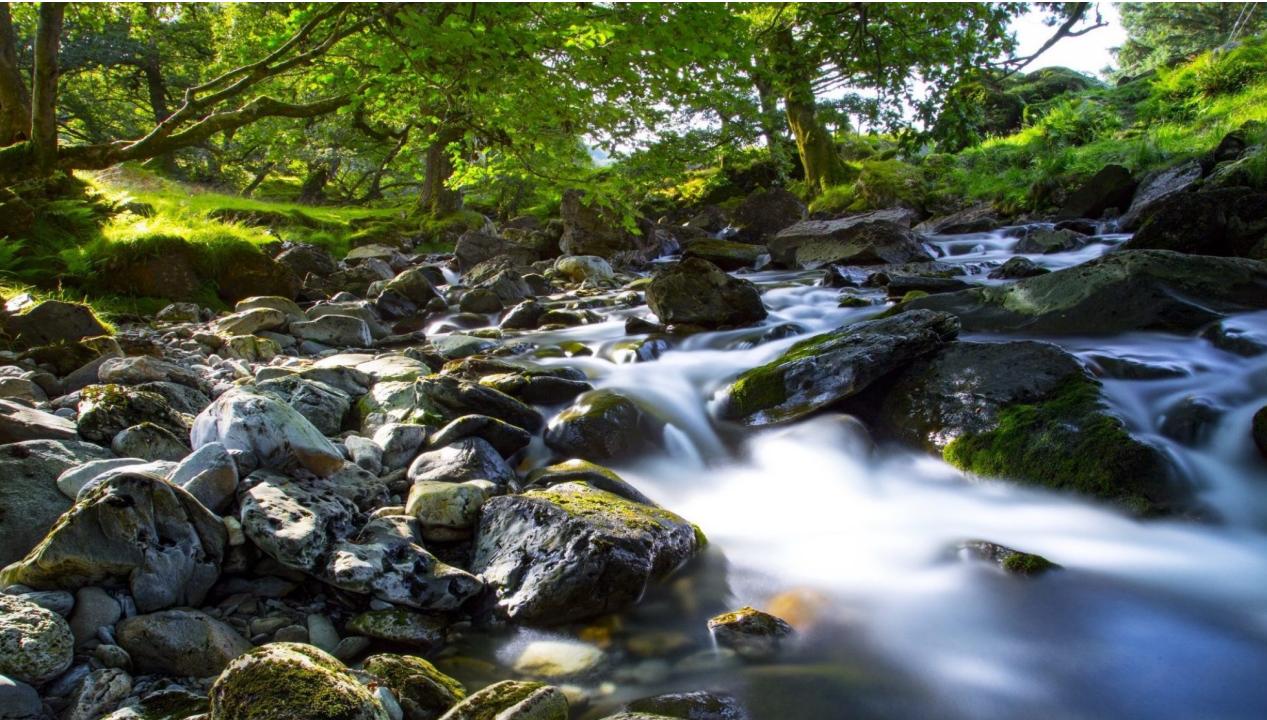
@manudellavalle - <http://emanueledellavalle.org>

**Continuous massive flows
than you cannot stop**

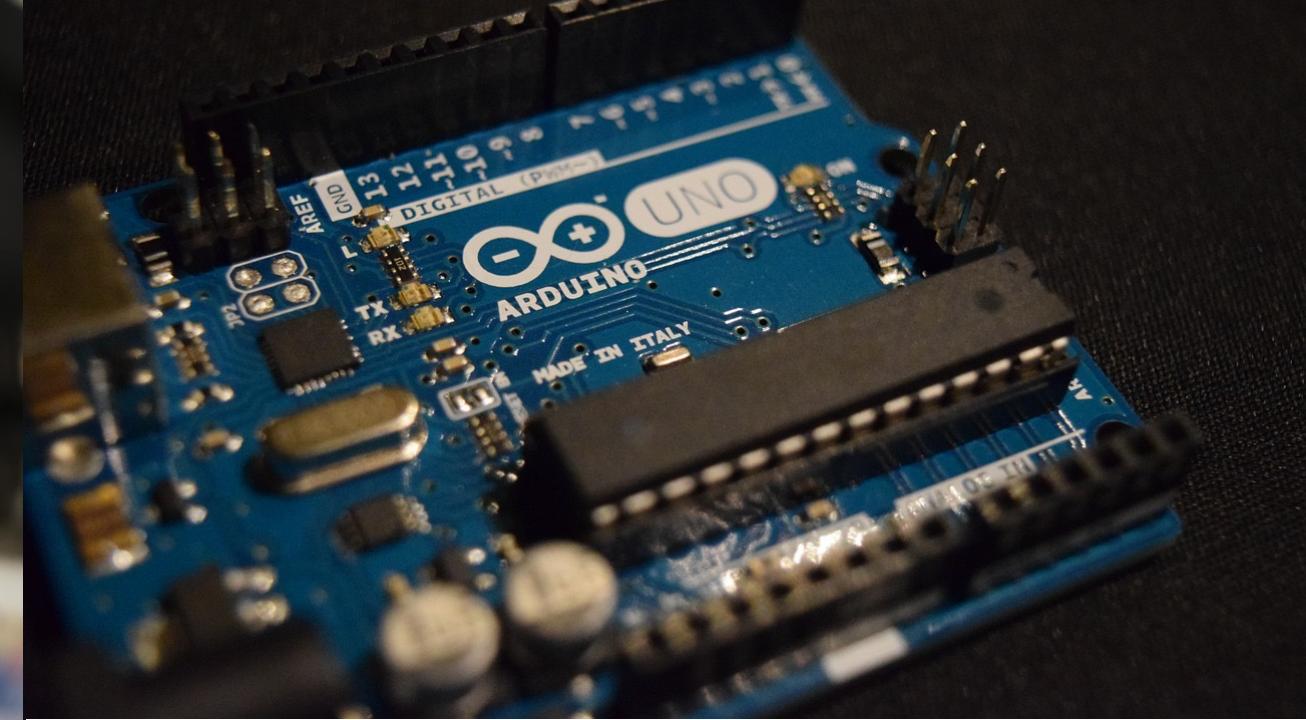
Image by Bruno /Germany from [Pixabay](#)

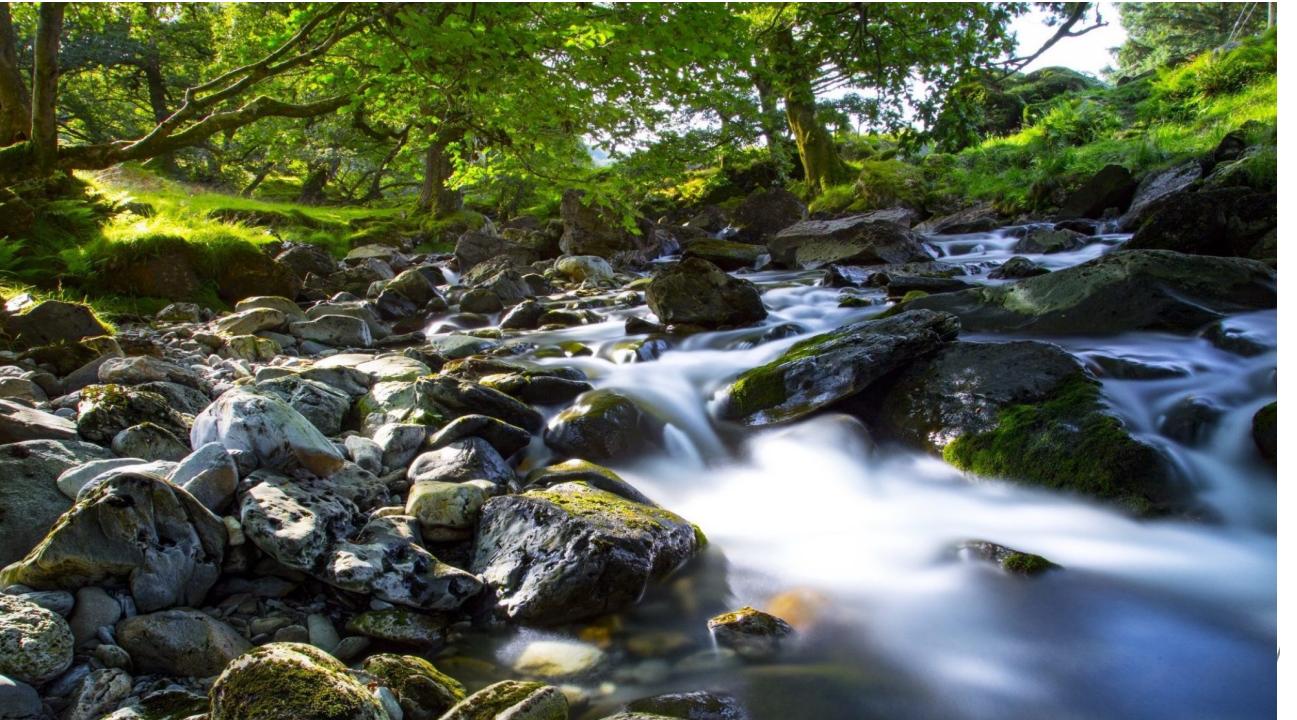


e.g., physical or cyber
alarms

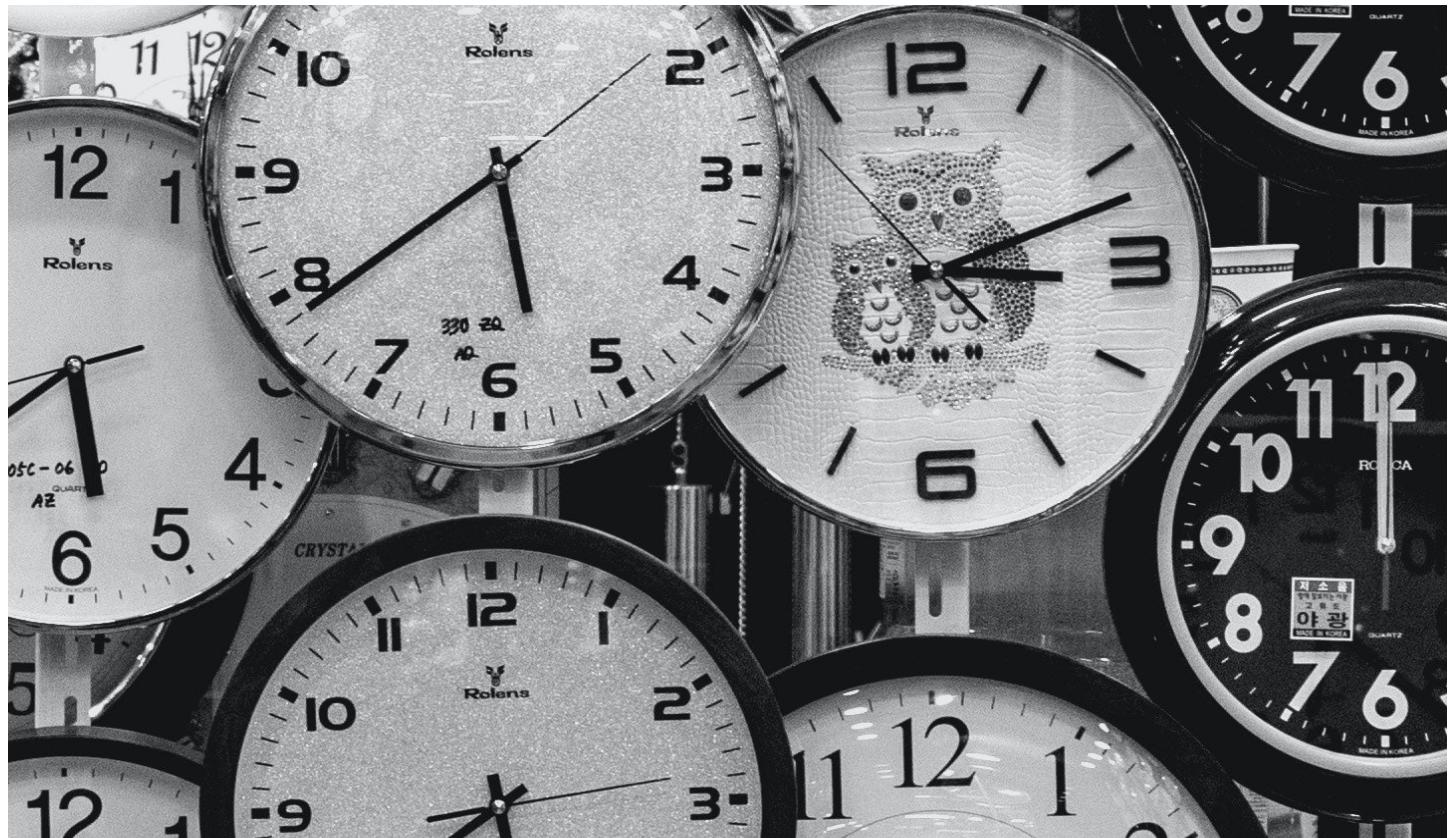


**Continuous numerous flows
that can turn into a torrent**

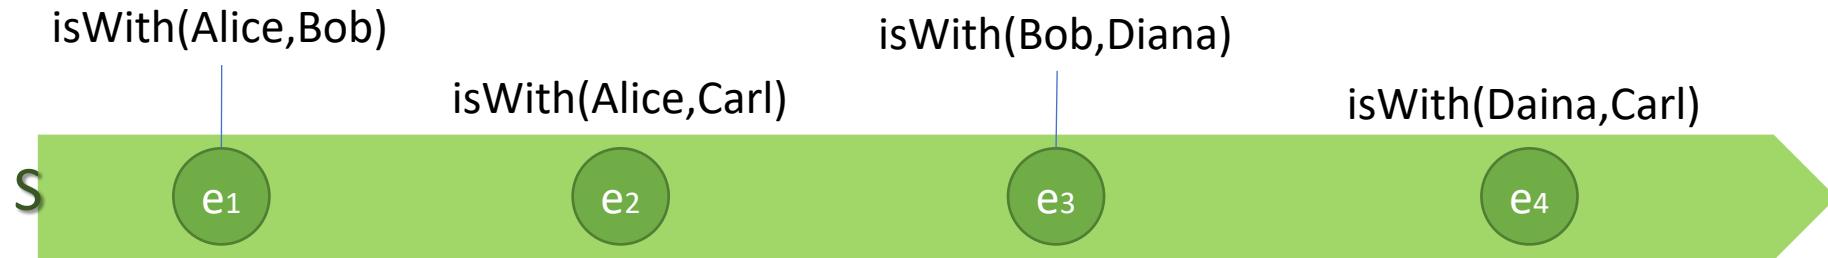




2nd premises: time (models) matters!

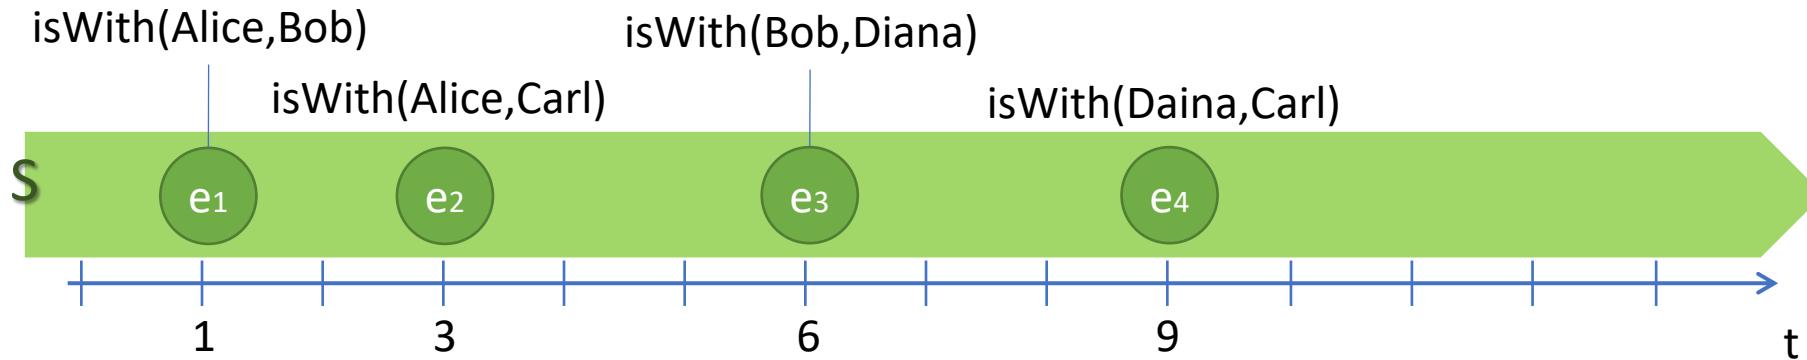


Stream-only Time Model



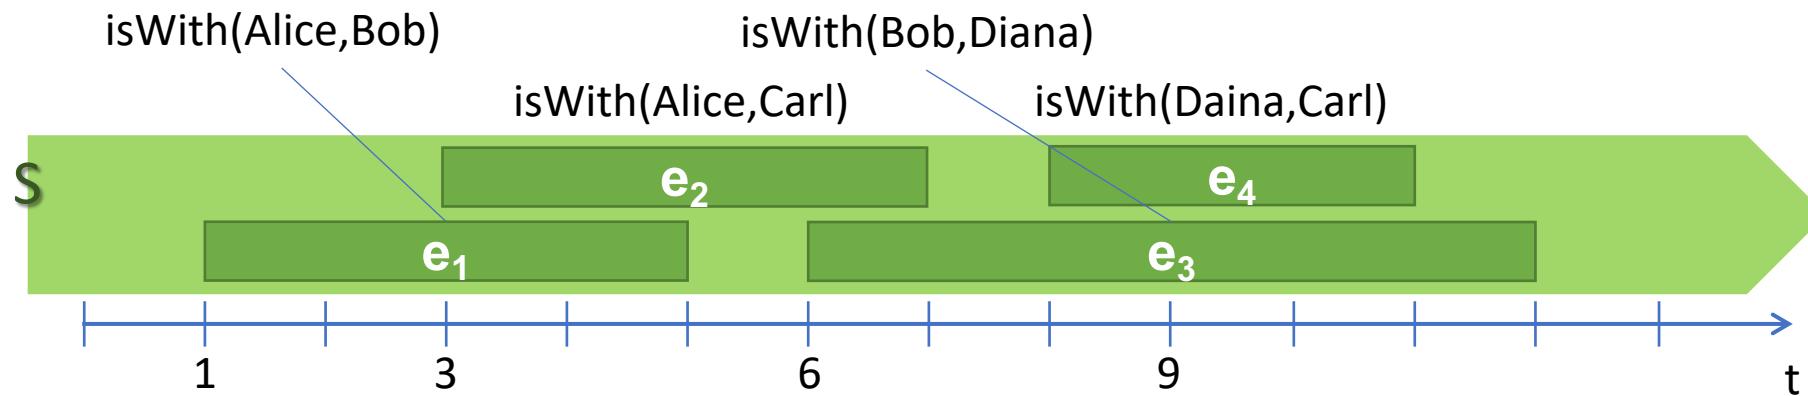
- The order can be exploited to perform queries
 - Does Alice meet Bob **before** Carl?
 - Who does Carl meet **first**?

Absolute Time Model



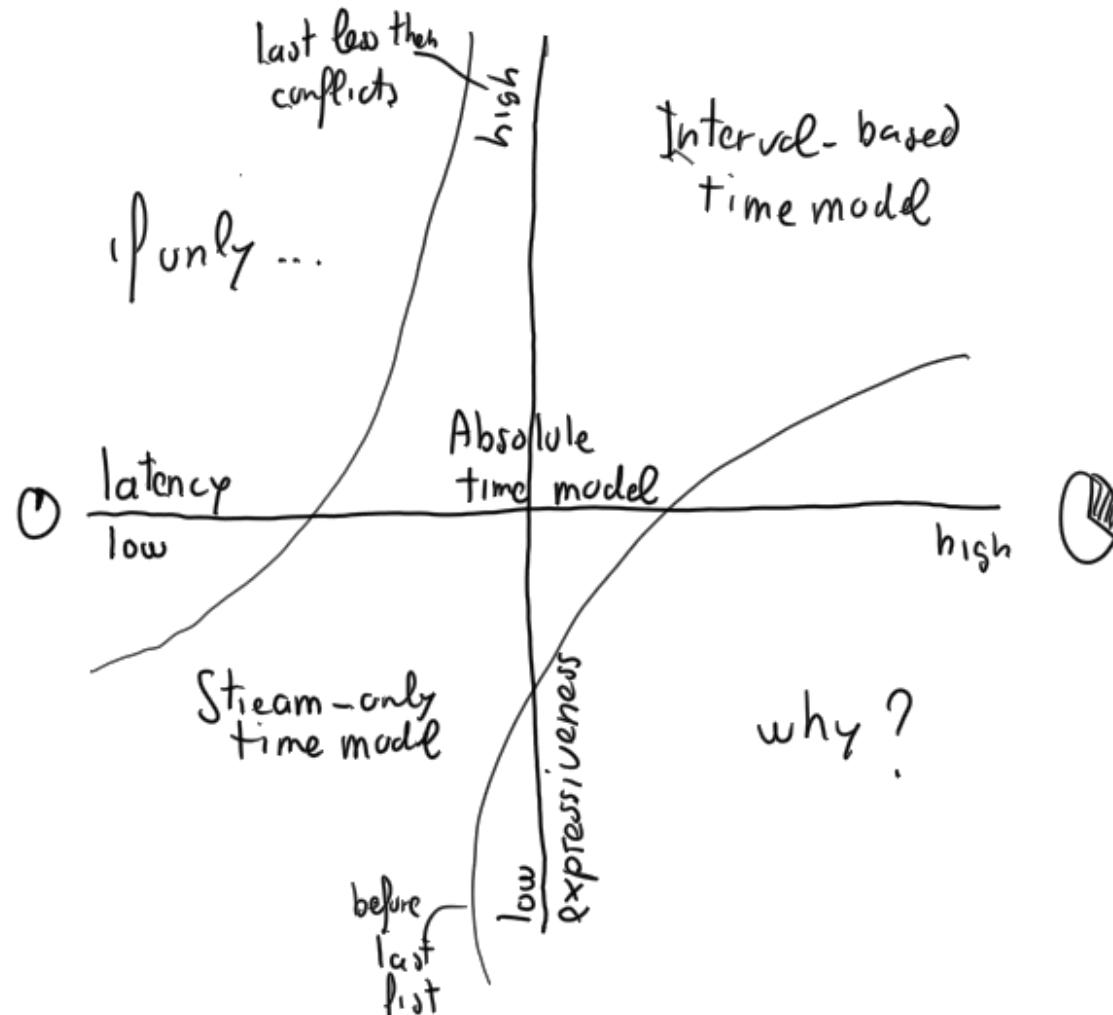
- We can ask the queries in the previous slide
- We can start to compose queries taking into account the time
 - How many people has Alice met **in the last 5m?**
 - Does Diana meet Bob and then Carl **within 5m?**

Interval-base Time Model



- We can ask the queries in the previous two slides
- It is possible to write even more complex queries:
 - Which are the meetings the **last less than 5m**?
 - Which are the meetings with **conflicts**?

A first trade-off: latency vs. expressiveness



Can we tame data velocity?

A long-exposure photograph of a waterfall cascading down a mossy rock face into a pool of water. The water appears as a bright, blurred stream against the dark, textured rock and lush green foliage.

Taming
*Myriads of tiny flows
that you can collect*
with
Event-based Systems

On Events and decoupling

Events

an immutable and append-only stream of “business facts”

Decoupling

means decentralizing the freedom to act, adapt, and change

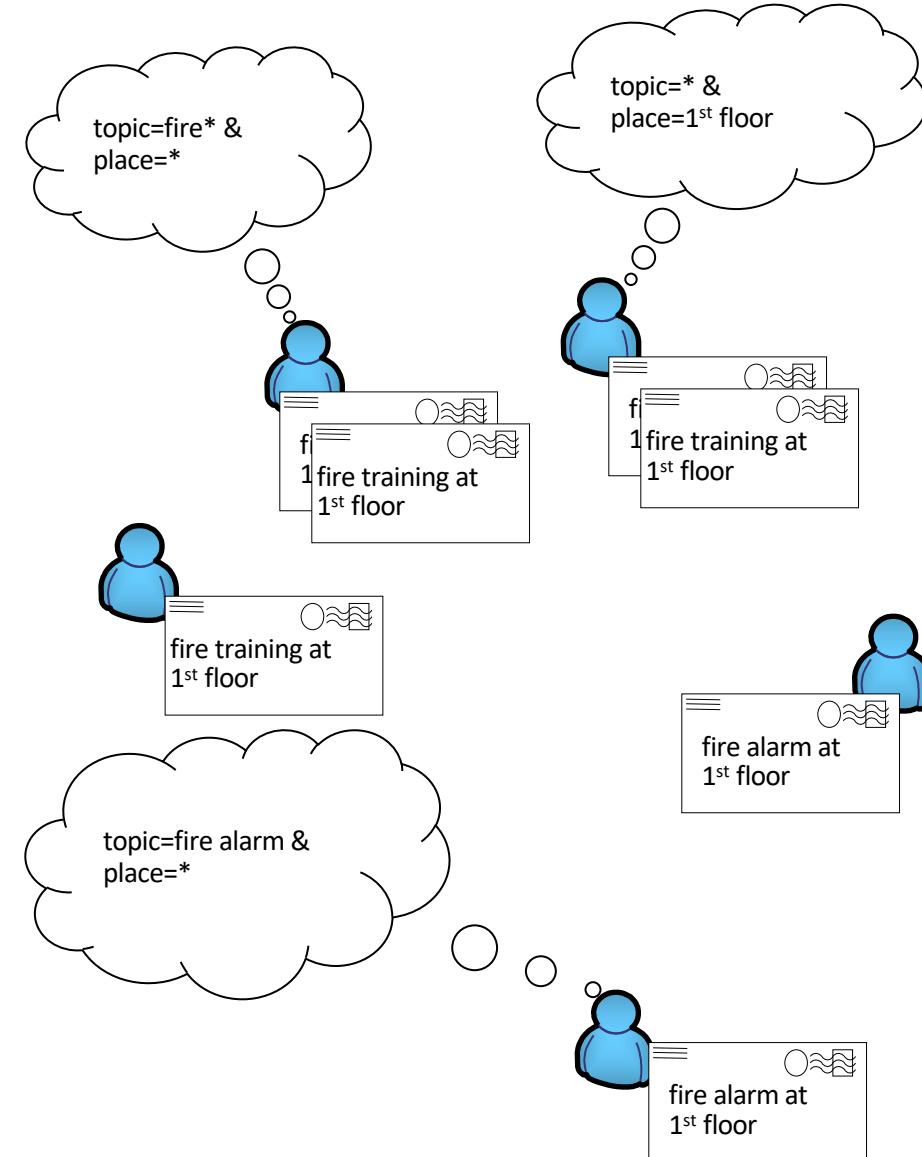
Event-based systems

- Events as “first-class” entities
- Services emit and consume events asynchronously
- Better decoupling with respect to APIs (no temporal coupling)

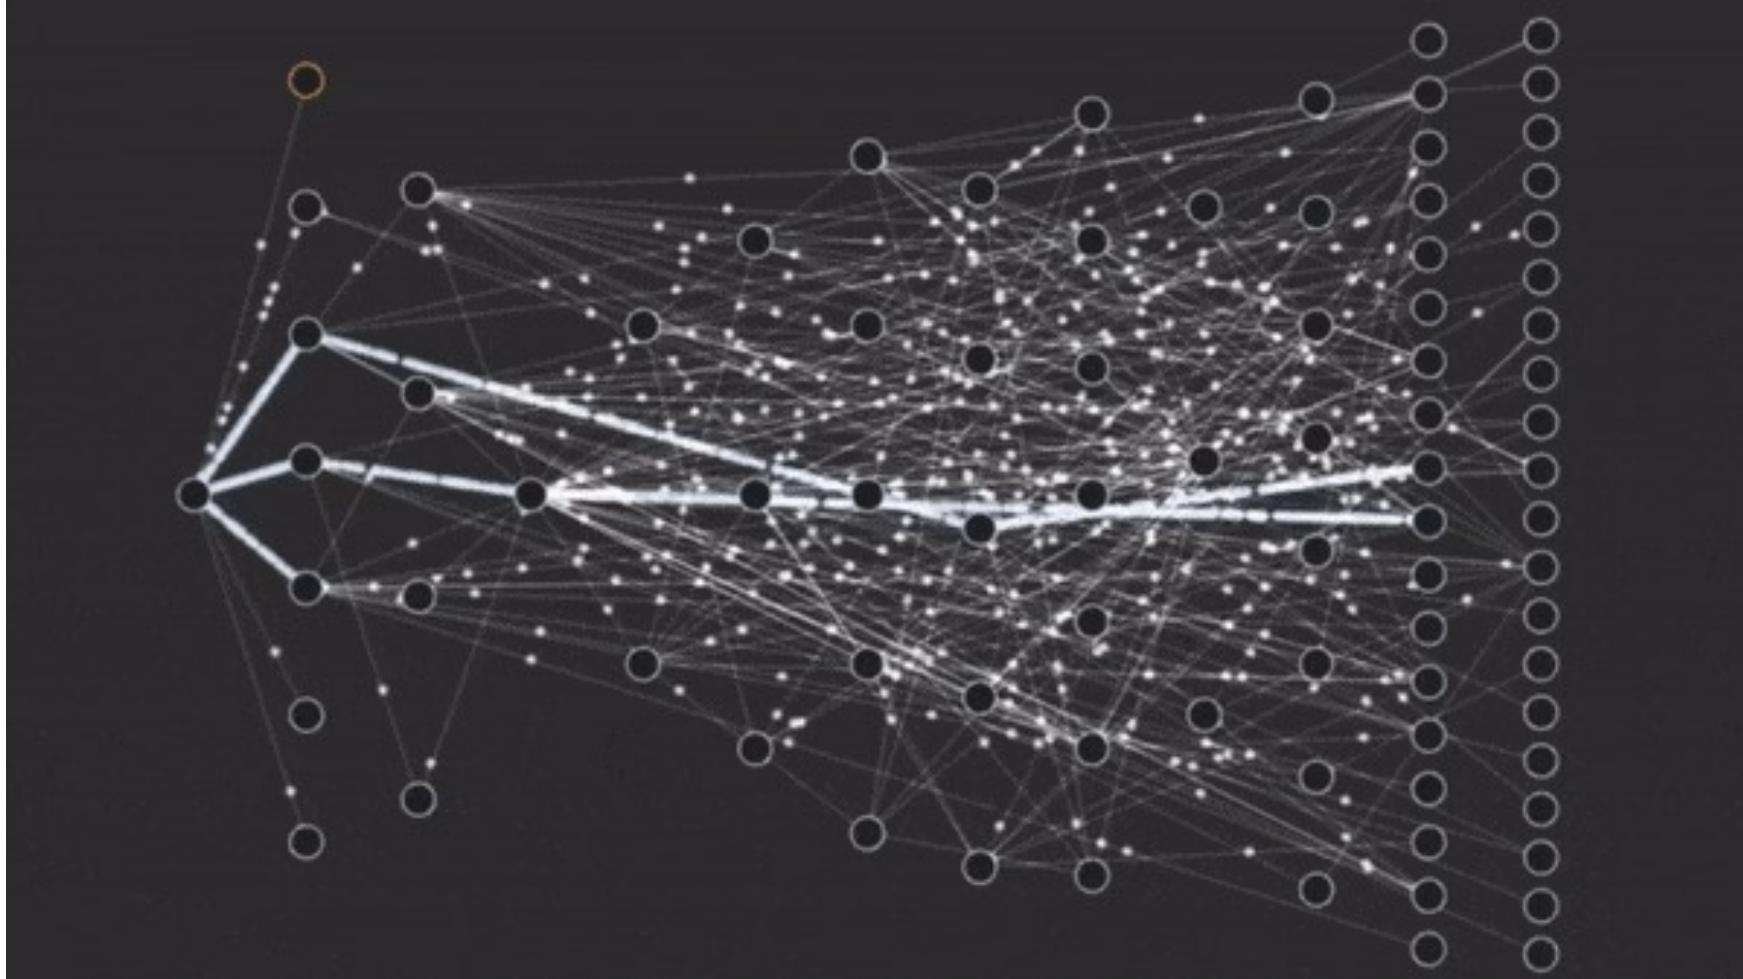
Eugster, P. T., Felber, P. A., Guerraoui, R., & Kermarrec, A. M. (2003). **The many faces of publish/subscribe.** *ACM computing surveys (CSUR)*, 35(2), 114-131.

Publish/subscribe

- Components
 - collaborate by exchanging *events*
 - *publish events they observe*
 - *subscribe to the events they are interested in*
- Communication is:
 - Purely message based
 - Asynchronous
 - Multicast
 - Implicit
 - Anonymous



Event-driven architecture @ Netflix



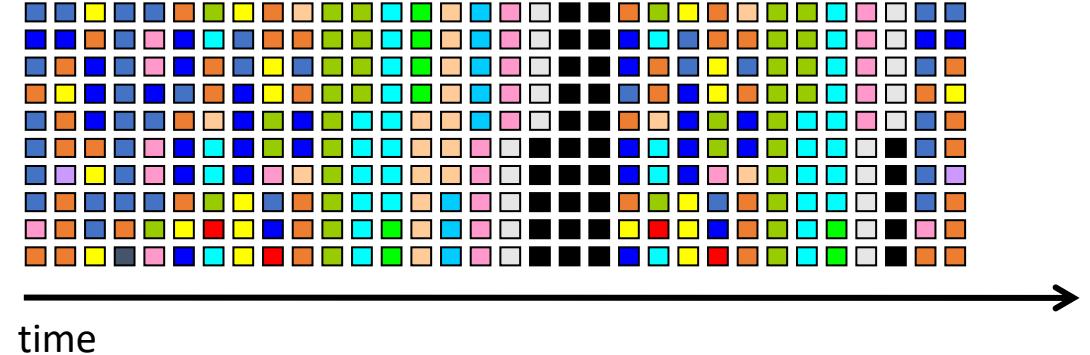
[src. <https://labs.sogeti.com/microservices-architectures-stay/>]

The background image shows a wide river flowing through a city. A small, densely forested island is located in the middle of the river. On the island, there is a small wooden platform or walkway where several people are standing. The water on the left side of the image is a clear, dark greenish-blue, while the water on the right side is a muddy, light brown color, indicating a mix of different water sources or perhaps a dam upstream. In the far background, buildings and trees are visible along the riverbank.

Taming
Continuous massive flows
than you cannot stop
with
Data Stream Mng. Systems

Data Stream Management Systems (DSMS)

- Data streams are (**unbounded**) sequences of time-varying data elements
- Represent:
 - an (almost) “continuous” flow of information
 - with the recent information being more relevant as it describes the current state of a dynamic system



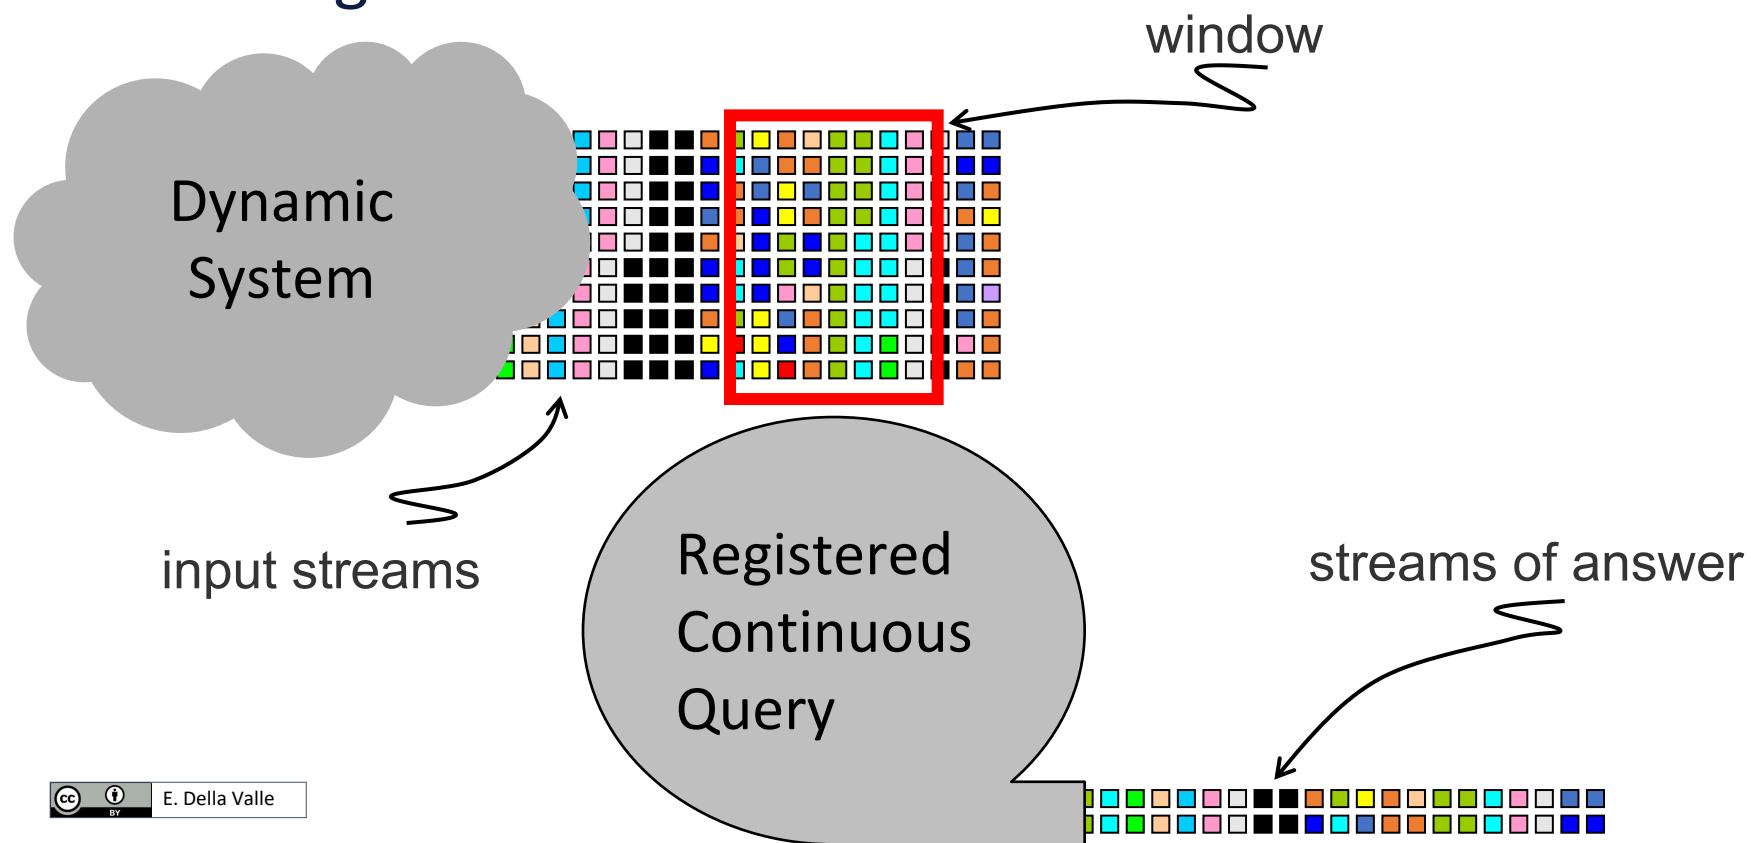
Data Stream Management Systems (DSMS)

- The nature of streams requires a paradigmatic change*
 - **from persistent data**
 - one time semantics
 - **to transient data**
 - continuous semantics

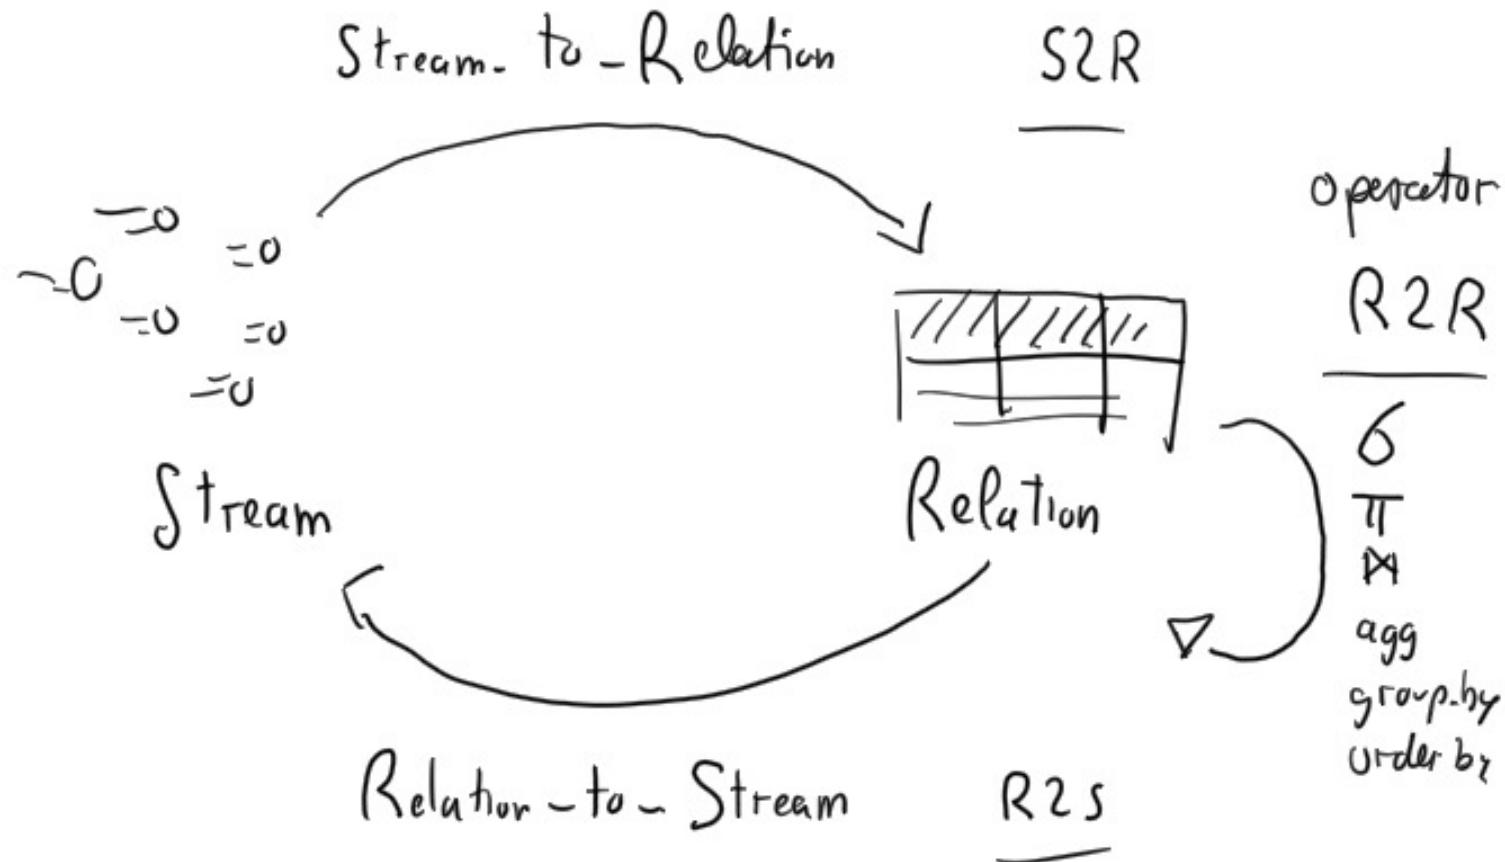
* This paradigmatic change first arose in DB community in **the late '90s**

Continuous Semantics

- Continuous **queries registered** over streams that are observed through **windows**



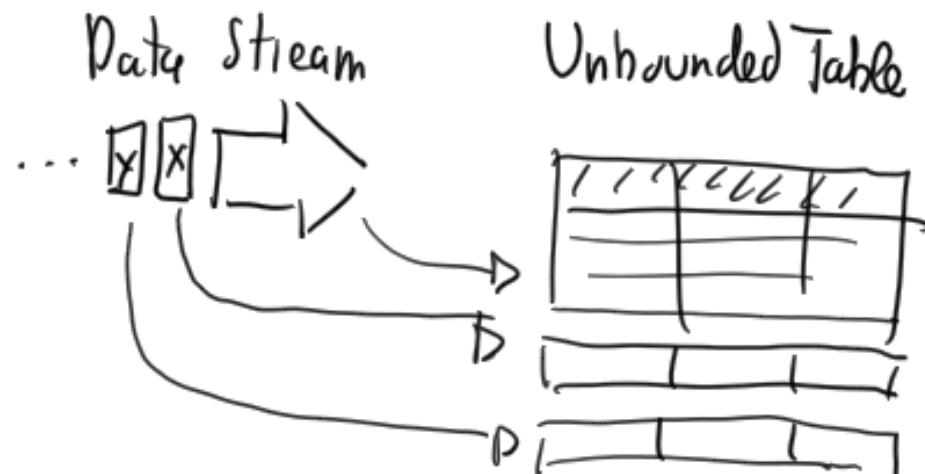
DSMS Semantics



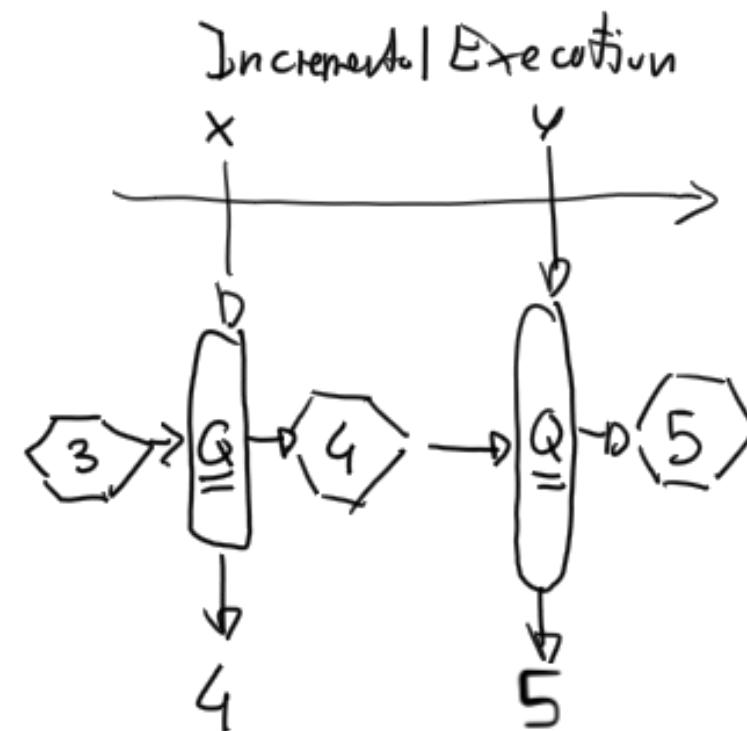
Arasu, A., Babu, S. and Widom, J., 2006. **The CQL continuous query language: semantic foundations and query execution.** *The VLDB Journal*, 15(2), pp.121-142.

Modern DSMS

Logic View



Physical View



$Q = \text{count}$

SQL Extensions

Windowed aggregation

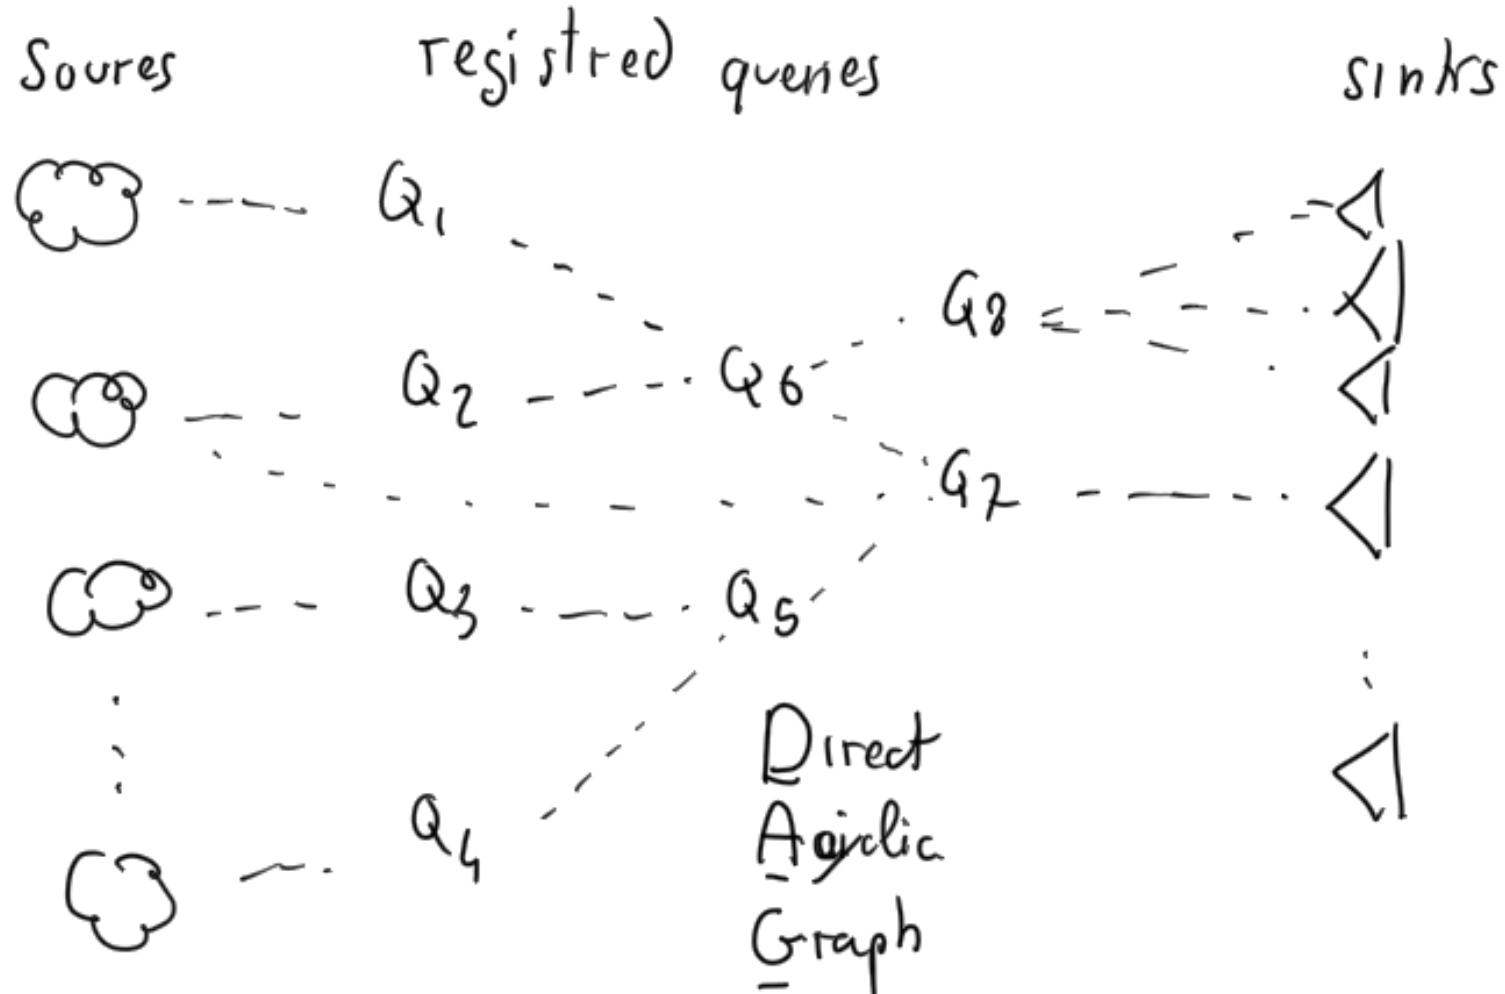
```
SELECT x, COUNT(*), SUM(y)  
FROM s1 WINDOW  
      TUMBLING (SIZE 1 MINUTE)  
GROUP BY x  
EMIT CHANGES;
```



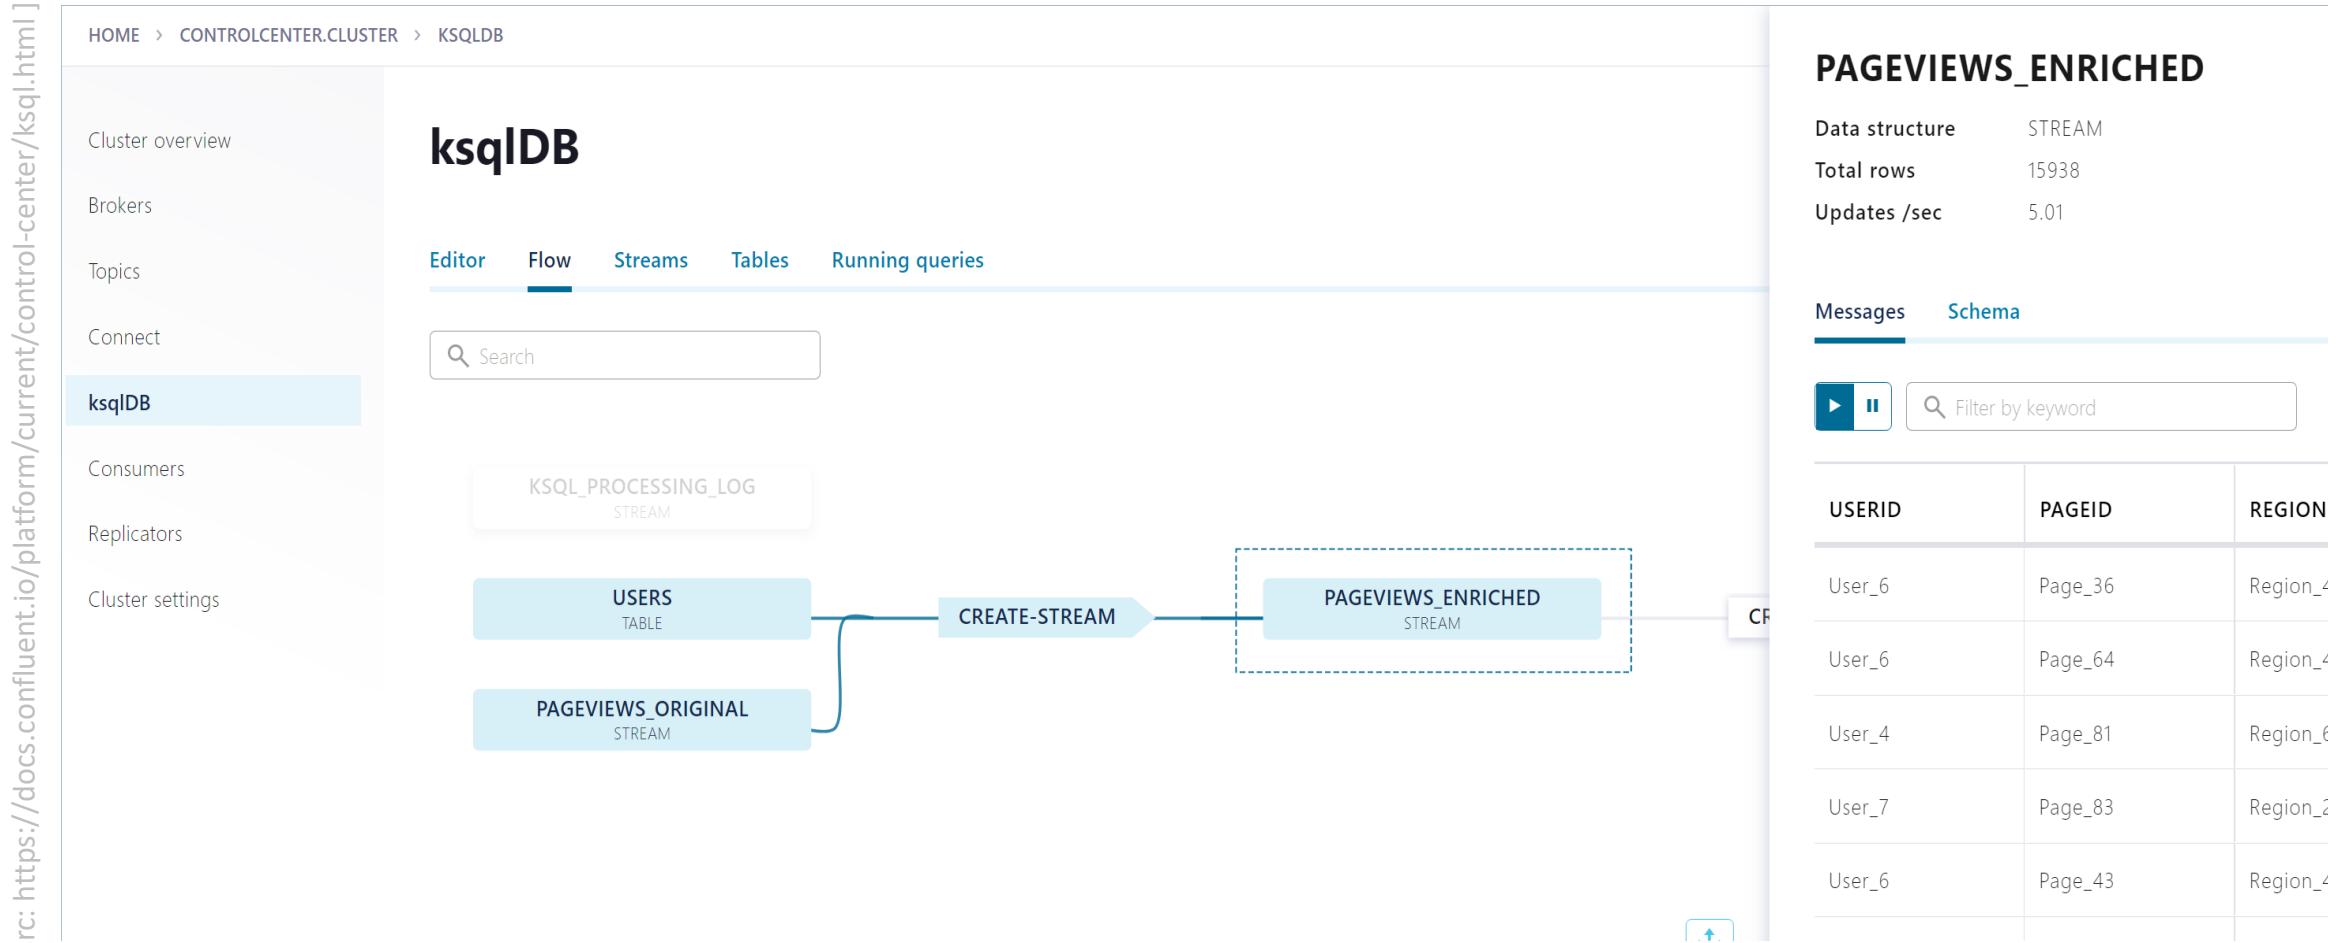
Join two streams together

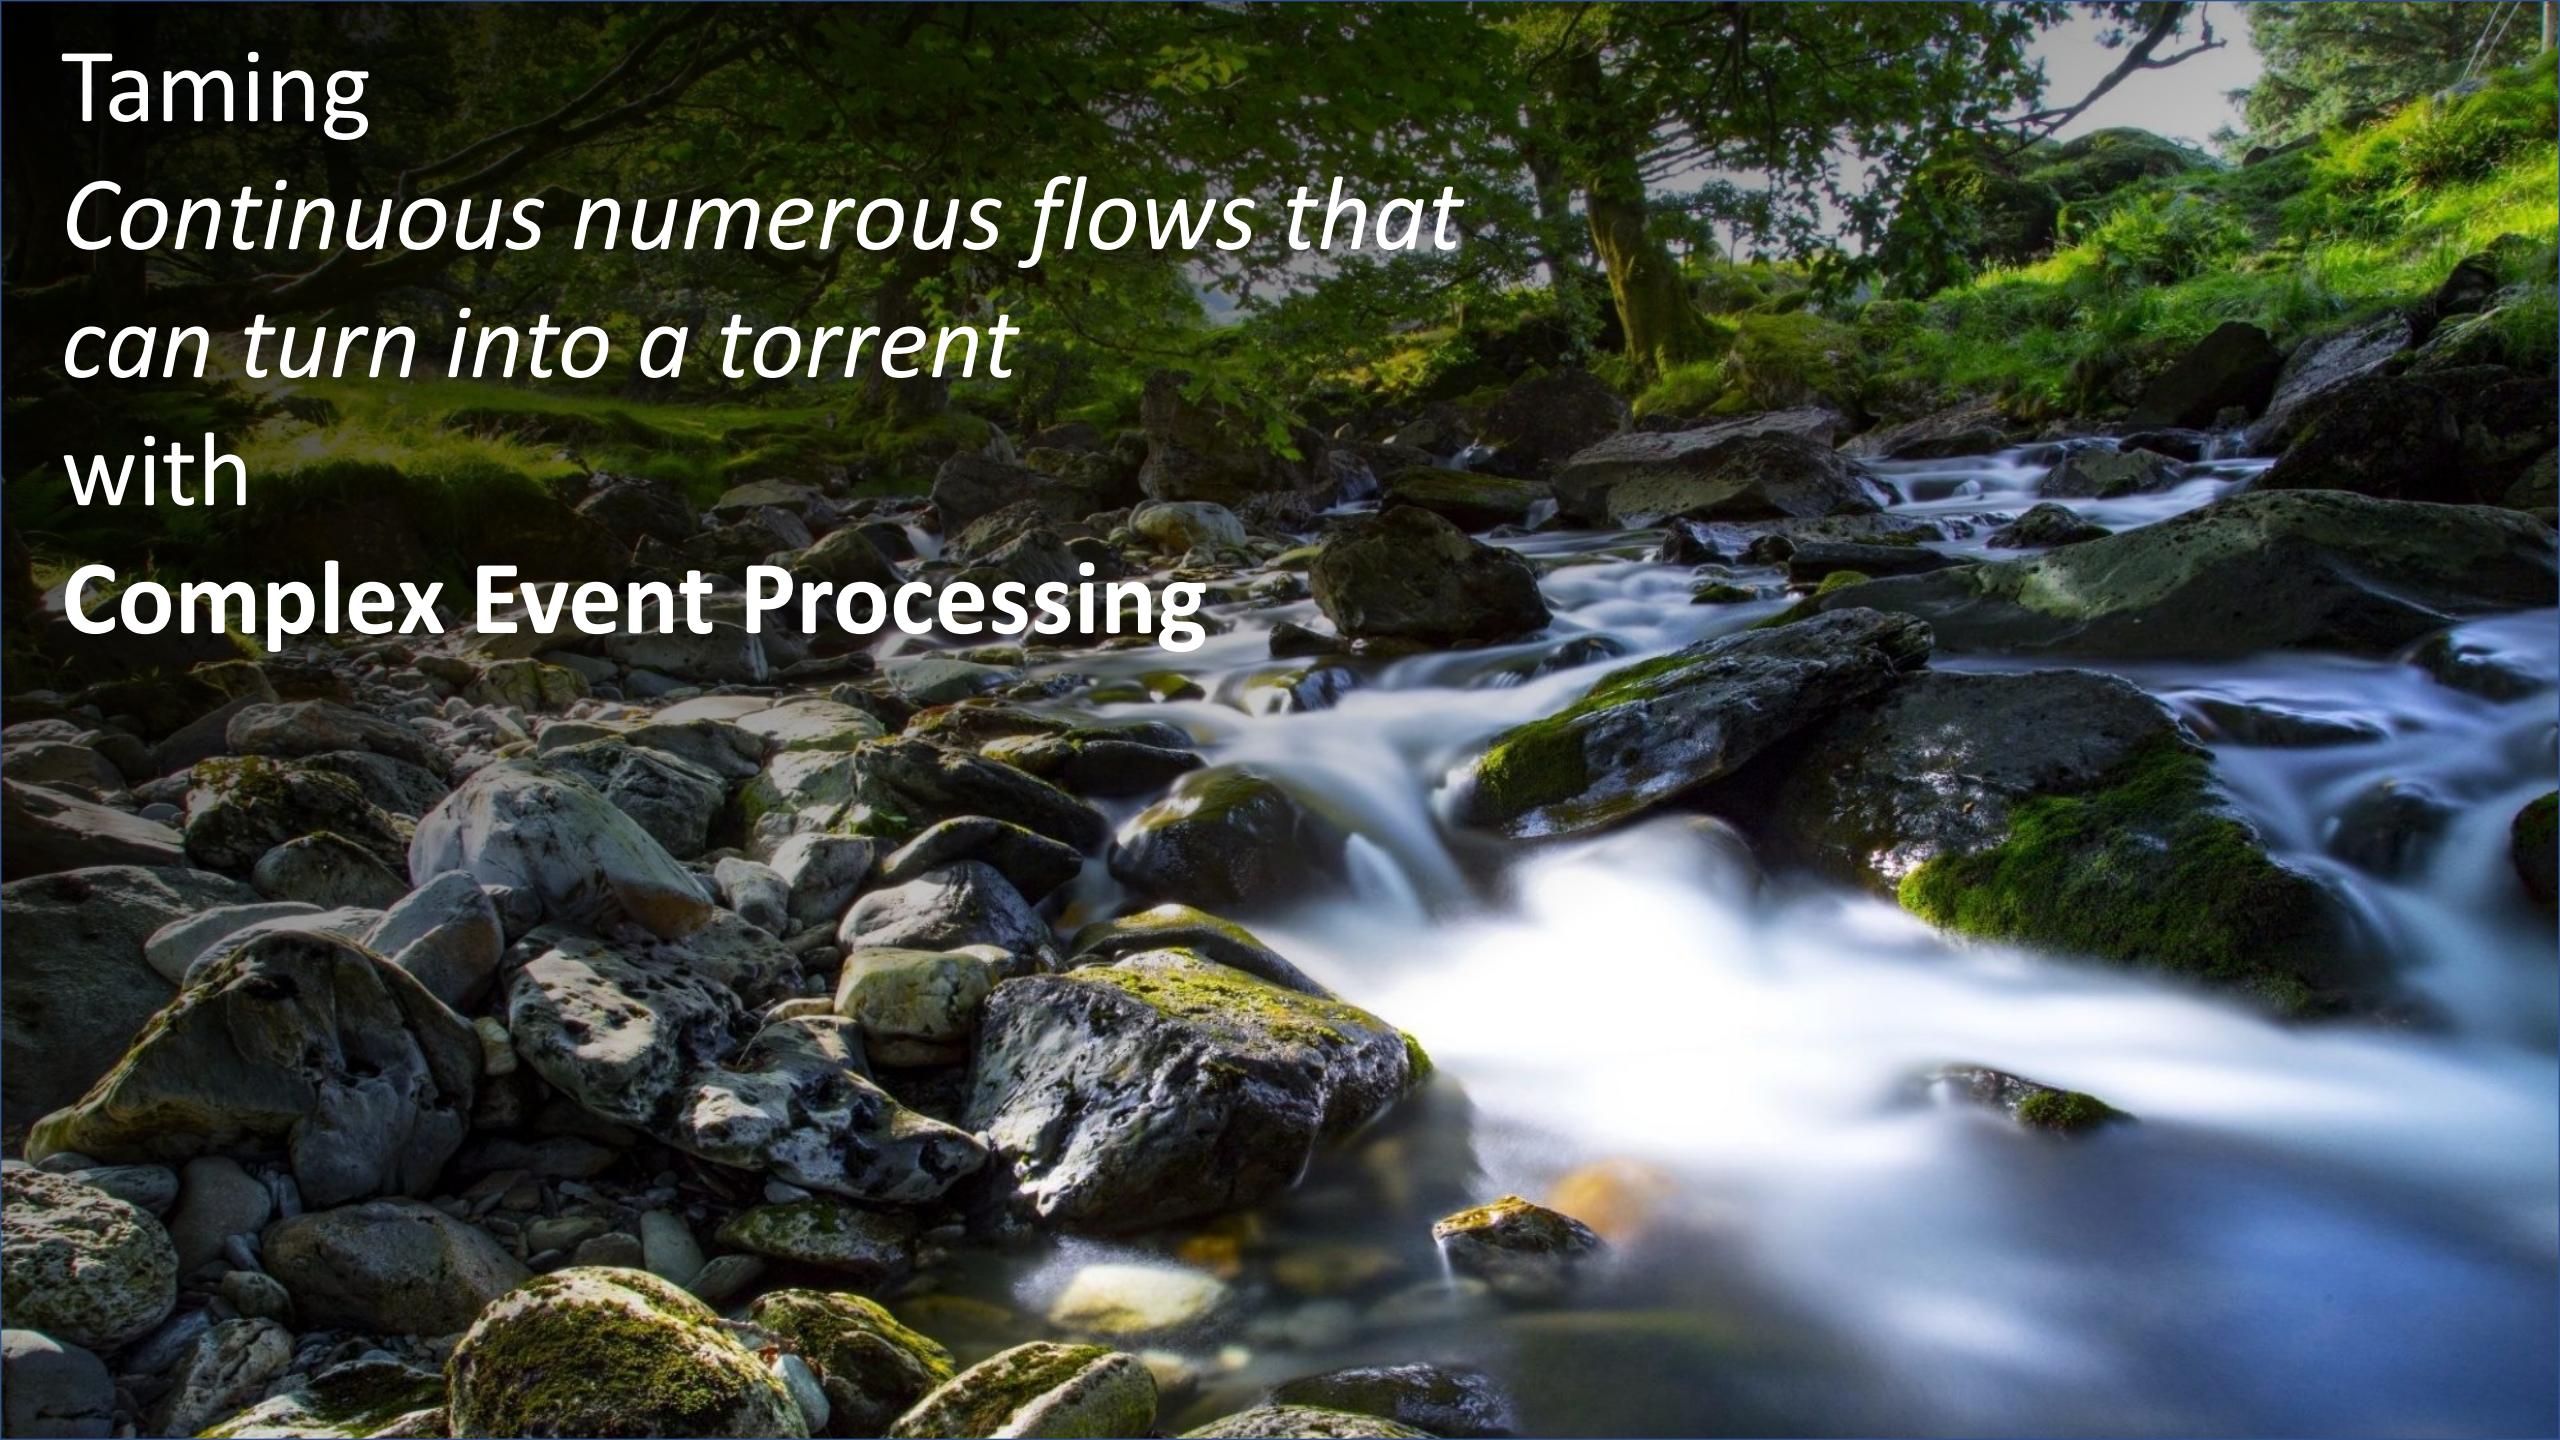
```
CREATE STREAM s3 AS  
SELECT s1.c1, s2.c2  
FROM s1 JOIN s2  
      WITHIN 5 MINUTES  
ON s1.c1 = s2.c1  
EMIT CHANGES;
```

Typically queries are placed in a network



Typically queries are placed in a network

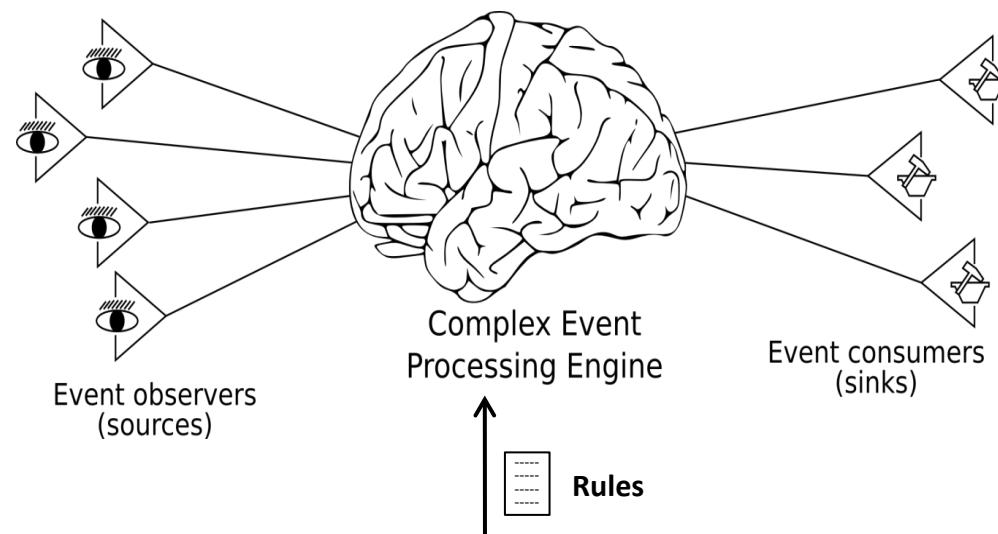




Taming
*Continuous numerous flows that
can turn into a torrent*
with
Complex Event Processing

Complex Event Processing (CEP)

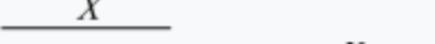
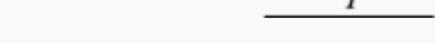
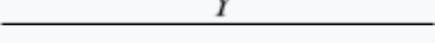
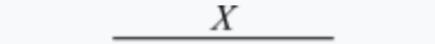
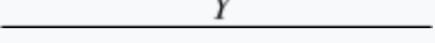
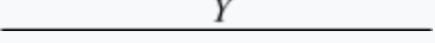
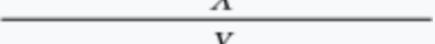
- CEP systems adds the ability to deploy *rules* that describe how composite events can be generated from primitive (or composite) ones
- Typical CEP rules search for *sequences of events*
 - Raise C if A→B
- Time is a key aspect in CEP



Cugola, G., & Margara, A. (2012). **Processing flows of information: From data stream to complex event processing**. *ACM Computing Surveys (CSUR)*, 44(3), 1-62.

CEP semantics

- subsets of Allen's Interval Algebra

Relation	Illustration	Interpretation
$X < Y$		
$Y > X$		X takes place before Y
$X \mathbf{m} Y$		
$Y \mathbf{mi} X$		X meets Y (<i>i</i> stands for <i>inverse</i>)
$X \mathbf{o} Y$		
$Y \mathbf{oi} X$		X overlaps with Y
$X \mathbf{s} Y$		
$Y \mathbf{si} X$		X starts Y
$X \mathbf{d} Y$		
$Y \mathbf{di} X$		X during Y
$X \mathbf{f} Y$		
$Y \mathbf{fi} X$		X finishes Y
$X = Y$		X is equal to Y

[src: https://en.wikipedia.org/wiki/Allen%27s_interval_algebra]

Detecting Languages



- Event Processing Language (EPL)

```
insert into alertIBM
select *
from pattern [
    every (
        StockTick(name="IBM", price>100)
        ->
        (StockTick(name="IBM", price<100)
            where timer:within(60 seconds))
    );
]
```

Write a query which alerts on each IBM stock tick with a price greater than 100 followed by an IBM stock tick lower than 100 within 60 seconds

- Supported since 2006 Esper by EsperTech Inc. and Oracle CEP

More Expressive Detecting Languages

```
1 stream<MatchT> Matches = MatchRegex(Quotes) {  
2     param  
3         pattern : ". rise+ drop+ rise+ drop* deep";  
4         partitionBy : symbol;  
5         predicates : {  
6             rise = price>First(price) && price>=Last(price),  
7             drop = price>=First(price) && price<Last(price),  
8             deep = price<First(price) && price<Last(price) };  
9     output  
10        Matches : symbol=symbol, seqNum=First(seqNum),  
11                         count=Count(), maxPrice=Max(price);  
12 }
```

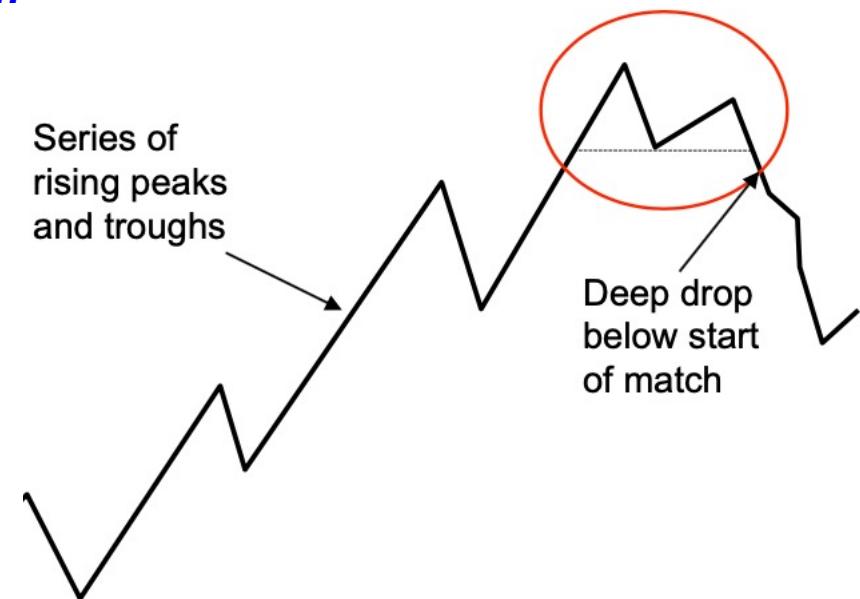
Composite events

Simple events

Regular expression

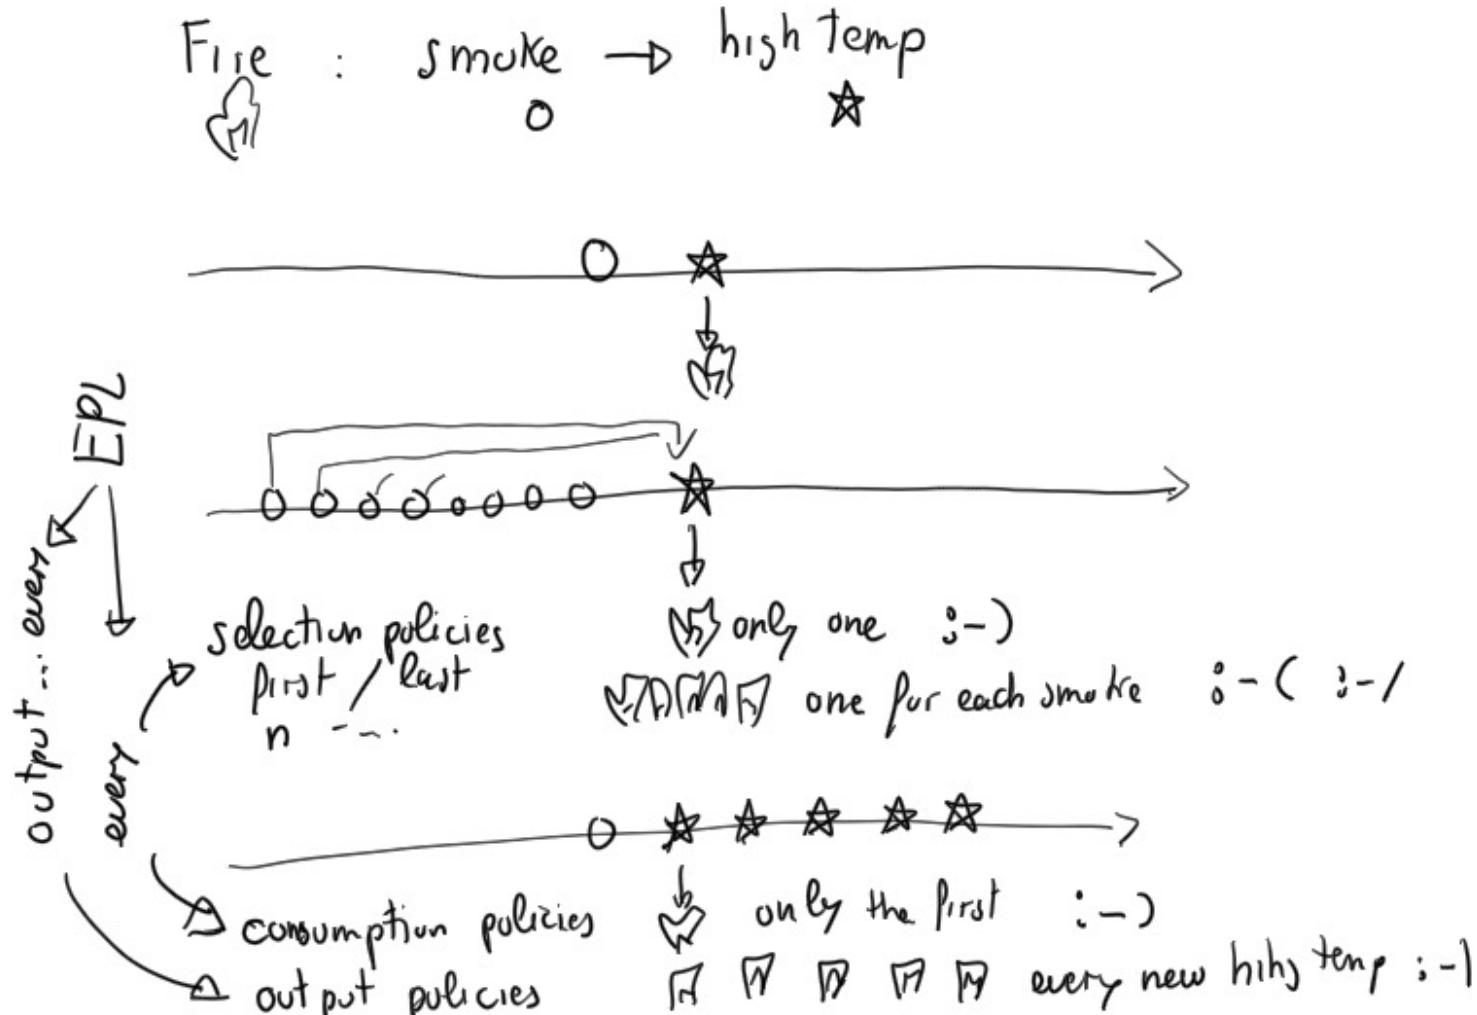
Key

Aggregation

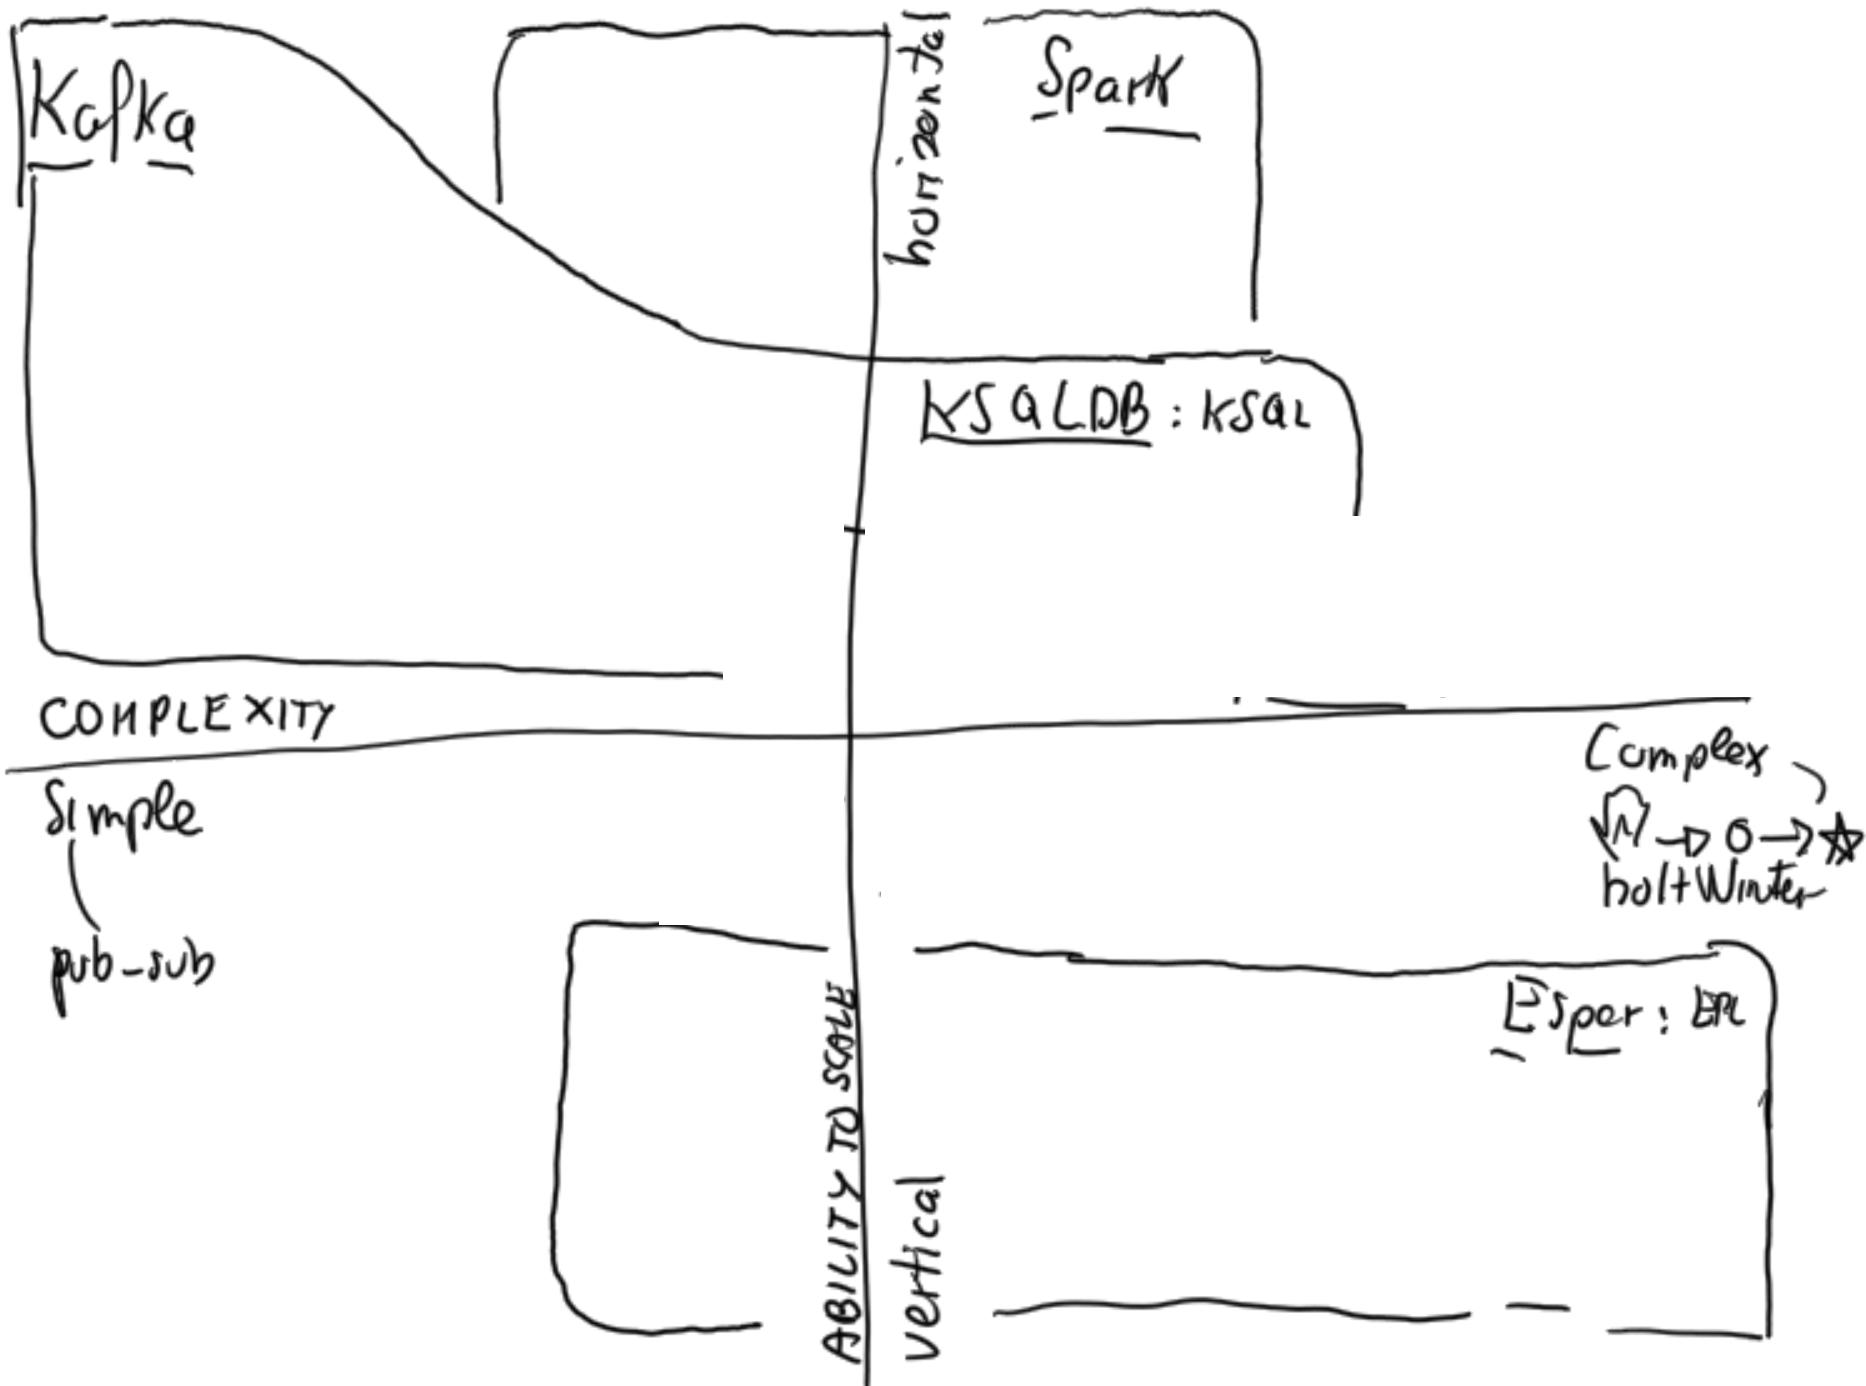


→ **Operator only, no extensions to SPL syntax**

Why can event-based systems fight the torrent effect?



Wrapping up



Thank you for your attention!

Questions?

Credits

- These slides are partially based on
 - "A modeling framework for DSMS and CEP" by G. Cugola and A. Margara presented in the PhD course on "Stream and complex event processing" offered by Politecnico di Milano in 2015.
 - <http://www.streamreasoning.org/courses/scep2015>
 - "Tutorial: Stream Processing Languages" by Martin Hirzel, IBM Research AI. 1 November 2017 Dagstuhl Seminar on Big Stream Processing Systems
 - <https://materials.dagstuhl.de/files/17/17441/17441.MartinHirzel1.Slides.pdf>
 - "The Death and Rebirth of the Event Driven Architecture" by Jay Kreps, Confluent, Kafka Summit 2018
 - <https://www.confluent.io/kafka-summit-london18/keynote-the-death-and-rebirth-of-the-event-driven-architecture>
 - "Time Series Databases" by Dmitry Namiot
 - <https://www.slideshare.net/coldbeans/on-timeseries-databases>
 - "Fractal Tree Indexes : From Theory to Practice" by Tim Callaghan
 - <https://www.slideshare.net/tmcallaghan/20131112-plukfractaltreestheorytopractice>

Taming Velocity

A Tale of Four Streams

Emanuele Della Valle

prof @ Politecnico di Milano

founder @ Quantia Consulting